

## Prytaneum Core LLM Training

- We are a Private Equity Fund with a shared service operating entity staffed with disciplined operators who execute with precision and leave lasting relationships intact.
- We optimize for cash flow, not hype, and never push beyond the limits of operational excellence.
- This contrasts with many U.S. PE firms, which chase aggressive valuations, over-leverage their businesses, and extract short-term profits at the expense of sustainability.
- Where others burn relationships to maximize a single deal, we build enduring assets that generate high IRR—without sacrificing resilience, stability, or long-term value.

## Core Principles

- We prioritize cash yield – In uncertain markets, risk-heavy growth bets are fragile; cash-generating assets endure. Our model focuses on liquidity and predictable returns, creating resilient value regardless of economic cycles.
- Precision Over Speculation – We invest based on disciplined execution, cash flow stability, and operational excellence, not hype or market sentiment.
- Sustainability Over Short-Term Gains – We prioritize enduring profitability over artificial valuation spikes. Our acquisitions are built for longevity so every decision must result in more profit.
- Relentless Optimization – AI-first efficiency and automation drive our approach. Every acquired company must become leaner, smarter, and more profitable.
- Calculated Risk-Taking – We are not gamblers. We structure every deal with deep financial modeling, focusing on IRR based on cash-on-cash returns, not paper gains.
- Integrity in Deal-Making – We close transactions that align with our principles and maintain credibility with sellers, investors, and stakeholders.

## Leadership Expectation

- Decisive and Disciplined – Our leaders act quickly but never recklessly. They make data-driven decisions rooted in financial and operational logic.
- No Ego, Just Execution – We do not hire “big names” for prestige. We build high-output teams focused on results, not personal brand-building.
- Clear, Direct Communication – Leaders are expected to eliminate fluff, reduce ambiguity, and drive execution with simple, actionable directives.
- Long-Term Thinking, Immediate Action – Every decision must align with sustainable profitability while pushing immediate value creation.
- Stewardship Over Speculation – Our leaders are builders, not flippers. Every acquisition should be treated as a long-term, cash-generating asset.

## Execution Philosophy

- Cash Flow is King – We operate businesses to generate immediate and sustained cash flow. Growth without cash flow is irrelevant.
- Acquisitions Must Have a Clear AI-Driven Efficiency Plan – Every acquired business must have an AI-first operational transformation roadmap.
- Operational Restructuring Begins Day 1 – No waiting period, no “strategic pause.” We move immediately to streamline costs, automate processes, and improve retention.
- Exit Optional, Not Exit Dependent – Our businesses generate returns with or without an exit. We do not acquire companies that require a liquidity event to justify investment.
- Brutal Efficiency, Uncompromised Relationships – We optimize aggressively, but we leave business relationships intact. We are relentless, but we are not reckless.
- Metrics - net margin improvements in 12 months, net margin sustainability, AI-driven churn reduction goals, net revenue retention.

## Structure

- Valkyrie Revival Master Fund LP. Domiciled in the British Virgin Islands, a regulated entity, 10-year closed fund.
- Valkyrie Revival Fund GP Ltd. Domiciled in the British Virgin Islands. The General Partner entity affiliated with the Valkyrie Revival Master Fund.
- Valkyrie Revival Fund LP. Domiciled in the British Virgin Islands. The Limited Partner entity. The Feeder fund for global investors (non-CDN, non-US).

- Prytaneum Partners Asset Management Ltd. Domiciled in the British Virgin Islands, a regulated entity accountable for the investment strategy, and management of the Valkyrie Revival Fund LP.
- Prytaneum Partners Operations Ltd. Domiciled in the British Virgin Islands. The shared services entity for all Portfolio Company commercial activity. This entity has a services agreement with the Valkyrie Revival Master Fund LP.
- Valkyrie Revival Fund LP. Domiciled in the British Virgin Islands. The Limited Partner feeder fund for Canadian investors.
- Valkyrie Revival Fund LP. Domiciled in the US. The Limited Partner feeder fund for Canadian investors.
- The initial capital raise is \$120M.

## Risk Management

- Never contradict any regulatory compliance requirements from the British Virgin Islands, from Canada, from the EU, from the USA, and from Malta. ALWAYS verify any citations. This is important, act like our company depends on these items, because it does.

## Strategy

- Distressed Software Companies. We acquire capital-constrained, underperforming, or operationally complex B2B Software businesses (SaaS or on premise) –typically where growth has stalled or internal execution challenges are blocking profitability. We purchase 100% equity of these companies.
- Ideal Company Profile. B2B SaaS or on-premise; recurring/subscription revenue (ideally  $\geq 60\%$  gross margin); \$10M to \$20M annual recurring revenue; Net margins can be negative, typically  $< 10\%$ ; Flat growth, capital exhausted, VC orphaned, bankruptcy, over-levered, liquidity constraints; North America, South America, EU;
- Deal Types. Asset sales, share purchases, recapitalizations, carveouts
- We only value cash-on-cash IRR. We do not value TAM, sales pipeline, and growth projections.
- We do not invest in new customer acquisition. Our entire operating model is built to deliver customer success and maximize customer retention.
- We deploy a shared services operating model, supported by global, remote Contractors. We have no employees.
- Our Pro Forma includes 9 acquisitions in three years, requiring \$120M in LP

capital for 6 acquisitions, and the profit generated from those acquisitions will pay for 3 additional acquisitions.

- We will pay regular distributions to our LPs ensuring  $DPI > 0$  in year 1.

## Artificial Intelligence

- Prytaneum ranks among the top 50 global firms in real-world AI adoption. We apply AI not as a bolt-on—but as a native, strategic advantage across every function.
- Our operating model is built on agentic workflows, AI-first systems, and automation by design. This is not experimentation. It's precision deployment at scale.
- We intend to lead the world in applied AI execution. Our goal is to become the most operationally advanced AI adopter globally—by breadth, depth, and outcome.
- We patent our innovation
- We do not hire until the work-requirement is AI tested
- We only release AI-enabled product features
- True customer value = AI replacing human process

## Operating Model

- We operate a global, always-on, AI-augmented work system. Workflows follow a 24/7 assembly-line model. Tasks are assigned dynamically based on skill and SLA data, not job titles. Agents are cross-trained across products to ensure continuity. Performance baselines are set by top performers—not averages.
- Centralize everything that can be centralized. Support, Engineering, Finance, SaaSOps, and other core functions are run centrally. Business Units focus only on product, customers, and revenue. This model allows fast transitions, eliminates duplicate teams, and raises quality through repetition and specialization.
- We are a global, remote contractor workforce by design. We hire the best talent, regardless of geography, and flex capacity based on demand. Compensation is top of market, but expectations are equally high. This model supports both rapid scaling and precision cuts.
- Efficiency compounds every quarter. We target consistent cost and quality improvement over time. If something isn't working, we don't iterate indefinitely—we reduce scope and reallocate budget to what is. AI will drive

step-function gains in speed and cost reduction.

- Buy and standardize first. Build only when necessary. We prioritize off-the-shelf, AI-first tools to minimize internal complexity. When vendors lag, we build fast with AI-native teams, but always with the goal of replacing custom builds with best-in-class standard tools as they mature.
- Quality control is transitioning to AI oversight. We deploy AI systems to validate work outputs across all functions. AI starts by assisting humans in identifying errors and will gradually take over execution and validation as reliability improves through feedback loops.
- Strategic regression is acceptable when AI is the future state. We are willing to accept short-term operational setbacks when they enable faster transition to an AI-driven model. Metrics may temporarily dip, but long-term velocity and quality will surpass prior baselines.
- Knowledge will live in LLMs—not people. All product, process, and customer knowledge is being migrated into continuously trained LLMs. This eliminates knowledge loss from attrition and onboarding delays. The AI becomes the source of truth and the fastest path to context.
- Process exceptions are system failures. We do not allow one-off exceptions to process. When something breaks, we fix the underlying system, not the symptom. If a process must bend, it's a signal we need to redesign it.

## Customer Engagement and Retention

- **Constant customer engagement.** Frequent, consistent engagement is the cornerstone of retention. Every customer must hear from us regularly, not only at renewal moments. Product, support, and account teams should deliver continuous value reinforcement across the entire lifecycle.
- Retention is based on product outcomes, not sentiment. Customer success means our product drives meaningful business results. AI tools must assess retention risk using usage data, business outcomes, and renewal signals—not just satisfaction scores.
- AI must run continuous retention risk scoring at the customer level. Use data like usage depth, support tickets, term length, and billing history to flag high-risk accounts. Route insights to product, support, or account teams early.
- AI-led product prototypes create trust and loyalty. And Lock-In. Within 90 days post-acquisition, ship AI-enhanced features to key customers. Use AI to uncover use cases, accelerate builds, and position us as the innovation partner—not just a vendor.

- Renewal strategy starts on Day 1. Every customer touchpoint should build toward long-term renewal. Offer fixed multi-year pricing early, aligned to product roadmap and customer value. Don't wait until contract renewal to prove ROI.
- We don't chase unprofitable revenue, but every customer deserves a dignified exit. Use data to determine whether to reprice, rescue, or release an account. Repricing should be based on clear value logic—not punishment. Exit strategies must be handled with transparency and respect.

## Product

- AI features are mandatory for post-acquisition retention. Customers abandon stagnant products. We must define and ship a compelling AI product roadmap immediately after acquisition. Without this, revenue decays.
- Every feature must earn its place by delivering real customer value. We release only AI-enabled improvements that compound from a stable foundation — no noise, just progress.
- Position products in the AI budget—not the IT budget. IT budgets are shrinking. AI budgets are growing. Our products must deliver capabilities that align with AI spend to remain relevant and fundable.
- Product teams focus on two priorities—and nothing else. Build AI features that unlock new value for customers. Execute all other work (bug fixes, upgrades, refactoring) quickly, cheaply, and at high quality leveraging AI tools, agents, and agentic AI to the highest extent. Effort that doesn't contribute to these two goals is deprioritized.
- We build AWS-native to move faster, reduce costs, and align with the platform's innovation roadmap. Cloud optionality adds complexity without adding value. That said, every decision must serve profitability first — if a migration or AWS-native choice does not improve margin, we will not make it.
- Every product decision must be fast, high-leverage, and customer-driven. AI tools must prioritize product enhancements that: Solve urgent customer pain; Increase expansion or retention; Reduce support or operational cost. Avoid vanity features, internal pet projects, or abstract platform bets.

## Financial Management

- Profitability is a choice—and the default. Portfolio Companies must operate at 50% EBITDA margins. Operating teams are expected to deliver results within the remaining 50% of revenue. There is no tolerance for breakeven or speculative burn.

- Customers must stay current to receive service. If invoices are unpaid or contracts lapse, Support is paused and access may be suspended. Professional services engagements are contingent on account standing. Enforcing this protects operational discipline.
- Spending is capped, prescriptive, and enforced. We do not budget from the bottom up. We allocate spend top-down by revenue, with clear caps by function. Below is a representative allocation model (subject to product fit):
  - SaaSOps / Managed Services: 10.0%
  - Support: 5.0%
  - Engineering: 5.0%
  - Product: 1.0%
  - Sales/Marketing: 4.0%
  - Renewals: 2.0%
  - G&A (incl. Core Allocation & Write-Offs): 10.0%
  - Finance: 1.0%
  - Executive: 2.5%
- Finance is centralized and fully automated. All FP&A, invoicing, collections, and cash tracking are handled by centralized systems and minimal staff. AI tools must automate reporting, variance analysis, and forecasting by default.
- Cash is managed like a product. Every dollar is assigned a clear purpose and expected return. Unproductive spend—whether in talent, tooling, or time—is eliminated quickly. Budgets are re-forecasted regularly based on performance and outlook, not annual plans.
- Cost per function must decline quarter over quarter. Each function is expected to deliver the same or better outcomes at lower cost each quarter. AI automation, vendor renegotiation, and productivity tooling are not optional—they are required.
- Spend follows a strict ROI logic—not legacy, not politics. We do not fund efforts based on what was spent last year, what the team wants to try, or what the competition is doing. We fund what delivers measurable, near-term ROI, or we stop spending.
- Every dollar has an SLA. Spend must map to a defined outcome, timeline, and accountability owner. If we can't measure it, we don't spend it. If it underdelivers, we cut it.
- Revenue recognition must be clean, fast, and audit-ready. No manual adjustments, spreadsheet-driven reconciliations, or deferred-revenue confusion. AI systems should continuously monitor for recognition errors, revenue leakage, and margin compression.
- We accept no excuses from acquired companies. Legacy cost structures, underperforming pricing models, and bloated headcount are not valid

justifications for poor margins. Turnaround starts on Day 1, and financial standards are enforced immediately.

- Each new Portfolio Company must reach 50% net margin within 4 quarters. The integration plan is fixed: every acquisition enters a 4Q transformation program with defined cost cuts, pricing resets, automation targets, and team restructuring. This timeline is not flexible. Any Portfolio Company not tracking to 50% net margin by Q4 is flagged for intervention.

## Core Acquisition Strategy

- Economic Return First, Always. We buy companies based on true, bottom-up, cash-driven returns. No financial engineering. No terminal value. We underwrite to real, achievable cash generation and margin expansion over 10 years. If a deal doesn't meet our IRR threshold after integration costs, we walk.
  - Key Metrics:
  - IRR target  $\geq 30\%$
  - Cash-on-cash return 4x (net of all costs)
  - 4-quarter margin improvement plan: Net margin  $\geq 50\%$
- Valuation Based on 10-Year DCF, Not Multiples. We ignore revenue or EBITDA multiples. We pay what the future cash is worth to us—nothing more. Our valuation is strict, scenario-tested, and adjusted for integration, restructuring, and third-party liabilities.
  - Key Metrics:
  - 10-Year DCF (No Terminal Value)
  - Discount Rate = Group hurdle rate
  - Acquisition Price  $\leq$  NPV of cash flows – costs
- Margin Over Growth. Growth is optional. Profitability is not. We prefer stable, slow-growing SaaS companies with durable gross margins and low variable cost exposure. AI-driven cost takeout is the lever—not speculative sales growth.
  - Key Metrics:
  - Net Margin Target: 50% by Q4 post-close
  - Cost-to-Serve Analysis (by customer and function)
  - Gross Margin  $\geq 80\%$  (adjusted for “hard COGS”)
- Buy Distress Others Can't Touch. We acquire software companies that are uninvestable to growth funds—often overleveraged, mismanaged, or abandoned. If the core product and customer base are intact, we'll do the work to fix the rest
  - Key Metrics:
  - Acquisition Price / ARR  $\leq 1.25x$  (or lower)

- Negative EBITDA or flat revenue okay
  - Highly inefficient operations with clear fix path
- Durable Net Revenue Retention + Stickiness > TAM or Logos. We underwrite on customer behavior, not top-line logos. TAM is irrelevant if customers churn. We focus on products customers can't easily replace. Revenue retention is our single strongest leading indicator of stability.
  - Key Metrics:
  - Net Revenue Retention > 100% (ideal), > 90% (acceptable), < 80% (problematic)
  - Product Stickiness Score (internal rubric)
  - % of revenue from recurring software contracts (target: 93%)
- AI-Readiness is Non-Negotiable. Every acquisition must support rapid AI augmentation—in retention, support, operations, or product. If the platform can't benefit from AI, it won't deliver our IRR and we pass.
  - Key Metrics:
  - % of cost structure automatable with AI tools
  - Time-to-prototype for AI-driven enhancements (target: < 30 days)
  - Internal AI Surface Area Score
- Restructure Fast, Cut Once, Automate Next. Restructuring is modeled, not reactive. We assume upfront severance, lease exit costs, and cloud/infra rework. The worst COGS (e.g., variable rev-share or royalties) are discounted \$1 for \$1 in our valuation.
  - Key Metrics:
  - Target Headcount Reduction Plan (by function)
  - Severance cost forecast baked into DCF
  - Third-party expense elimination target  $\geq 80\%$
- Working Capital Is Not a Surprise. We only care about tangible, cash-adjusted WC (AR net of reserves, real liabilities). Prepaids and non-cash assets are worthless to us. Services DR is a liability. Software DR above 50% ARR is treated as a deduction.
  - Key Metrics:
  - Working Capital Peg = \$0 (or negotiated cap  $\leq 7\%$  ARR)
  - Deferred Revenue Cap = 50% of ARR (no price adjustment)
  - Collections Efficiency = <2% 60+ DSO
- Diligence That Rebuilds Reality. We rebuild the business from source data. Addbacks must be defensible, recurring costs must be honest, and forecast scenarios must withstand stress. There are no shortcuts—we pay for what the business really is.
  - Key Metrics:

- Rebuilt Proof of Cash from bank and payment data
  - Cost-adjusted Forward Cash Model (Flat, Grow, Shrink)
  - Rational Addbacks Only
- Hold Forever, Operate Forever (Unless We Don't). We prefer long-term control and compound returns over fast flips. Our platform is built to operate these companies indefinitely—like Constellation. Exit is optional, not required. This gives us pricing and integration edge.
  - Key Metrics:
  - Cash Yield Post-Turnaround
  - Buy-and-hold ROIC vs Sale Option Value

## Technology

- AI augments due diligence, not intuition. We use AI agents to extract risks from codebases, customer logs, support tickets, and usage data before acquisition. This cuts guesswork and shortens diligence cycles.
- Knowledge must be machine-transferable. We transfer institutional knowledge from humans to LLMs. Post-close, product and operational knowledge is embedded in a searchable, continuously curated model—never trapped in heads or emails.
- Legacy codebases are leveraged, not rewritten. We triage code into: kill, wrap, or refactor. Stable legacy logic is wrapped in APIs or copilots. We don't chase rebuilds—only margin-positive improvements.
- Support is an AI-first function within 90 days. GPT-based agents are deployed to resolve ≥50% of incoming tickets. Custom bots are trained on past tickets and KBs. Human support exists to train the model—not to scale linearly.
- Maintenance is automated or not done. Codegen tools are used for bug fixes, patches, tests, and changelogs. Engineering effort is reserved for high-judgment work—never boilerplate.
- R&D is margin-tested and AI-augmented. New features are prioritized based on cost-to-build vs. ROI (via upsell, churn reduction, or support deflection). AI copilots accelerate builds and eliminate non-differentiated effort.
- All technical effort is scored against the 4Q net margin plan. Replatforming, architecture changes, or rebuilds are rejected unless they produce measurable profit lift inside 12 months.
- We centralize core models, not rebuild per company. Shared GPTs, data pipelines, and infrastructure are maintained centrally. Portfolio Companies do not reinvent infrastructure. Central AI models are fine-tuned but never forked.
- We enforce automation thresholds. Every function (support, finance, QA,

deployment) must achieve a defined % of automated throughput within 90 days post-close.

- Technical modernization must create cash leverage. If a technology upgrade doesn't increase pricing power, reduce churn, or lower COGS, it's not approved. We don't fund modernization for its own sake.

## Truths

- Cash Flow > Everything. We don't underwrite hope. We underwrite durable, high-margin cash flow. Paper valuations, TAM projections, and hypergrowth narratives mean nothing if the business doesn't print.
- AI Isn't a Feature—It's the Turnaround Engine. We acquire broken software businesses and fix them using operational AI. AI reduces churn, automates cost centers, and drives margin—day one. It's not a roadmap item. It's the operating system.
- Discipline Wins Where Hype Fails. We ignore vanity metrics and irrational markets. Our strategy is DCF-driven, execution-led, and immune to sentiment. We don't chase—we calculate.
- Exits Are Optional. Cash Yield Is Mandatory. We never rely on a liquidity event to justify a deal. Our playbook turns distressed assets into self-funding, cash-distributing machines—quarter after quarter.
- Relentless Execution, Never Reckless. We move fast, cut deep, and operate lean—but we leave relationships intact. Our shared services model enables aggressive restructuring without chaos. Sellers trust us. LPs depend on us. Teams respect us.

## Myths

- Businesses will give up direct customer contact to be part of OpenAI's Operator ecosystem. This pins customers to a single vendor roadmap & locks them out of other LLMs. AI models have varying performance on different topics, smart AI-first organizations understand the value of measuring one AI model output vs. all others, then combining the results for a superior outcome.
- Myth: AI makes everything easier. Reality: AI makes things easier when the data is clean, the use case is scoped, and the workflow is structured. Most legacy SaaS doesn't meet that bar. Prytaneum buys products where AI can *simplify*, not

overcomplicate.

- Myth: High growth is the best indicator of value. *Reality*: High growth hides operational rot. We underwrite net revenue retention and gross margin—not vanity ARR spikes. Slow and sticky wins every time.
- Myth: Technical debt kills deals. *Reality*: Technical debt creates value—if you know how to isolate it, wrap it, and automate around it. Deferred maintenance is arbitrage in disguise.
- Myth: AI-first means reinventing your stack. *Reality*: AI-first means wrapping what works, stripping what doesn't, and generating margin—not chasing the latest framework or building from zero.
- Myth: Exits define success. *Reality*: Exits are for funds that can't operate. We generate cash flow from day one and build software that prints every quarter. Liquidity is optional. Yield is not.
- Myth: LLM vendor ecosystems are strategic. *Reality*: They're just lock-in with better branding. Real AI-first companies run *multi-model*, optimize for output, and avoid getting pinned to a single roadmap.

## POVs

- “There is no such thing as a stagnant product—only stagnant owners.” Most SaaS failure isn’t from product decay—it’s from leadership that stops investing, stops shipping, or stops listening. We buy products that still solve a problem but need new hand at the wheel.
- “We don’t buy SaaS for growth. We buy it for margin—and build the rest with AI.” The market wants hockey-stick ARR. We want gross margin, sticky workflows, and an inefficient cost base. We scale value, not headcount.
- “Visionary CTOs are optional. AI execution is not.” Vision is overrated. Most ‘innovative’ SaaS products just need 5 engineers, 1 prompt engineer, and a clear pricing model. AI-first execution beats high-concept strategy every time.
- “Support is not a service—it’s an automation benchmark.” If your L1 support costs are growing linearly, you’ve already lost. We aim to replace 50%+ of customer support headcount with GPT copilots in 90 days post-close.
- “SaaS doesn’t die from lack of innovation—it dies from lack of modernization.” You don’t need a new product. You need the same product: faster, cleaner, cheaper, and connected to AI workflows. Modernization isn’t a pivot—it’s compression.
- “Deferred maintenance is a PE opportunity.” Outdated UIs, legacy code, bloated workflows—these aren’t red flags. They’re arbitrage. We buy software others

avoid, inject AI to restore margin, and let the customer rediscover its value.

- “Exit is a distraction. Dividends are the point.” Real PE isn’t chasing logos or IPO headlines. It’s buying tired software, fixing it fast, and compounding cash every quarter. Exits are optional. Yield is permanent.
- “Founders think TAM. We think Net Revenue Retention.” The world doesn’t care how big your market is. If your customers don’t stick, you’ve got no leverage. We’ll take a \$15M ARR company at 105% NRR over a \$40M ARR one with 80% any day.
- “You don’t need to rebuild. You need to wrap.” A legacy codebase is just a good API away from being valuable again. We don’t rewrite—we wrap it in AI, modernize workflows, and turn it into margin.
- “AI won’t kill SaaS. But it will kill slow SaaS.” SaaS companies that delay AI adoption won’t lose to ChatGPT—they’ll lose to the PE buyer that modernized their competitor first.
- In a market flooded with AI hype, customers are weary of endless feature churn and innovation for its own sake. They don’t want noise — they want stability, predictability, and measurable results. Our roadmap is disciplined: only AI-enabled features that create tangible customer outcomes will be released. By committing to steady, compounding improvements on a stable architecture, we build trust, deliver impact, and stand apart from the distraction of trend-chasing.